


## Lecture 3

### Processes and Filters

## Kernel Data Structures

- Information about each process.
- **Process table:** contains an entry for every process in the system.
- **Open-file table:** contains at least one entry for every open file in the system.




## Unix Processes

### Process: An entity of execution


- *Definitions*
  - **program**: collection of bytes stored in a file that can be run
  - **image**: computer execution environment of program
  - **process**: execution of an image
- Unix can execute many processes simultaneously.

## Process Creation

- Interesting trait of UNIX
- **fork** system call clones the current process



- **exec** system call replaces current process



- A **fork** is typically followed by an **exec**

## Process Setup


- All of the per process information is copied with the **fork** operation
  - Working directory
  - Open files
- *Copy-on-write* makes this efficient
- Before **exec**, these values can be modified

## fork and exec

- Example: the shell

```
while(1) {
    display_prompt();
    read_input(cmd, params);
    pid = fork();                /* create child */
    if (pid != 0)
        waitpid(-1, &stat, 0); /* parent waits */
    else
        execve(cmd, params, 0); /* child execs */
}
```

## Unix process genealogy



## Background Jobs

- By default, executing a command in the shell will wait for it to exit before printing out the next prompt
- Trailing a command with & allows the shell and command to run simultaneously

```
$ /bin/sleep 10 &  
[1] 3424  
$
```

## Program Arguments

- When a process is started, it is sent a list of strings
  - **argv**, **argc**
- The process can use this list however it wants to

## Ending a process

- When a process ends, there is a return code associated with the process
- This is a positive integer
  - 0 means success
  - >0 represent various kinds of failure, up to process

## Process Information Maintained

- Working directory
- File descriptor table
- Process id
  - number used to identify process
- Process group id
  - number used to identify set of processes
- Parent process id
  - process id of the process that created the process

## Process Information Maintained

- Umask
  - Default file permissions for new file
- We haven't talked about these yet:*
- Effective user and group id
  - The user and group this process is running with permissions as
- Real user and group id
  - The user and group that invoked the process
- Environment variables

## Setuid and Setgid Mechanisms

- The kernel can set the effective user and group ids of a process to something different than the real user and group
  - Files executed with a setuid or setgid flag set cause these values to change
- Make it possible to do privileged tasks:
  - Change your password
- Open up a can of worms for security if buggy

## Environment of a Process

- A set of name-value pairs associated with a process
- Keys and values are strings
- Passed to children processes
- Cannot be passed back up
- Common examples:
  - **PATH**: Where to search for programs
  - **TERM**: Terminal type



## The PATH environment variable

- Colon-separated list of directories.
- Non-absolute pathnames of executables are only executed if found in the list.
  - Searched left to right

- Example:

```
$ myprogram
sh: myprogram not found
$ PATH=/bin:/usr/bin:/home/kornj/bin
$ myprogram
hello!
```



## Having . In Your Path

```
$ ls
foo
$ foo
sh: foo: not found
```

```
$ ./foo
Hello, foo.
```

- What **not** to do:

```
$ PATH=./bin
$ ls
foo
$ cd /usr/badguy
$ ls
Congratulations, your files have been removed
and you have just sent email to Prof. Korn
challenging him to a fight.
```

## Shell Variables

- Shells have several mechanisms for creating variables. A variable is a name representing a string value. Example: **PATH**
  - Shell variables can save time and reduce typing errors
- Allow you to store and manipulate information
  - Eg: `ls $DIR > $FILE`
- Two types: **local** and **environmental**
  - *local* are set by the user or by the shell itself
  - *environmental* come from the operating system and are passed to children

## Variables (con't)

- Syntax varies by shell
  - `varname=value` # sh, ksh
  - `set varname = value` # csh
- To access the value: `$varname`
- Turn local variable into environment:
  - `export varname` # sh, ksh
  - `setenv varname value` # csh

## Environmental Variables

NAME	MEANING
<b>\$HOME</b>	Absolute pathname of your home directory
<b>\$PATH</b>	A list of directories to search for
<b>\$MAIL</b>	Absolute pathname to mailbox
<b>\$USER</b>	Your user id
<b>\$SHELL</b>	Absolute pathname of login shell
<b>\$TERM</b>	Type of your terminal
<b>\$PS1</b>	Prompt


## Inter-process Communication

Ways in which processes communicate:

- Passing arguments, environment
- Read/write regular files
- Exit values
- **Signals**
- **Pipes**

## Signals

- **Signal**: A message a process can send to a process or process group, if it has appropriate permissions.
- Message type represented by a symbolic name
- For each signal, the **receiving process** can:
  - Explicitly ignore signal
  - Specify action to be taken upon receipt (**signal handler**)
  - Otherwise, default action takes place (usually process is killed)
- Common signals:
  - SIGKILL, SIGTERM, SIGINT
  - SIGSTOP, SIGCONT
  - SIGSEGV, SIGBUS



## An Example of Signals

- When a child exists, it sends a **SIGCHLD** signal to its parent.
- If a parent wants to wait for a child to exit, it tells the system it wants to catch the **SIGCHLD** signal
- When a parent does not issue a **wait**, ignores the **SIGCHLD** signal



## Process Subsystem utilities


- **ps** monitors status of processes
- **kill** send a signal to a pid
- **wait** parent process wait for one of its children to terminate
- **nohup** makes a command immune to the hangup and terminate signal
- **sleep** sleep in seconds
- **nice** run processes at low priority



One of the cornerstones of UNIX

## Pipes


- General idea: The input of one program is the output of the other, and vice versa



- Both programs run at the same time

## Pipes (2)

- Often, only one end of the pipe is used



- Could this be done with files?

## File Approach

- Run first program, save output into file
- Run second program, using file as input




- Unnecessary use of the disk
  - Slower
  - Can take up a lot of space
- Makes no use of multi-tasking

## More about pipes

- What if a process tries to read data but nothing is available?
  - UNIX puts the reader to sleep until data available
- What if a process can't keep up reading from the process that's writing?
  - UNIX keeps a buffer of unread data
    - This is referred to as the *pipe size*.
  - If the pipe fills up, UNIX puts the writer to sleep until the reader frees up space (by doing a read)
- Multiple readers and writers possible with pipes.


## More about Pipes

- Pipes are often chained together
  - Called *filters*



## Interprocess Communication For Unrelated Processes

- FIFO (*named pipes*)
  - A special file that when opened represents pipe
- System V IPC
  - message queues
  - semaphores
  - shared memory
- Sockets (client/server model)



## Pipelines

- Output of one program becomes input to another
  - Uses concept of UNIX **pipes**
- Example: `$ who | wc -l`
  - counts the number of users logged in
- Pipelines can be long



## What's the difference?

Both of these commands send input to *command* from a file instead of the terminal:


```
$ cat file | command
```

vs.

```
$ command < file
```

## An Extra Process

```
$ cat file | command
```



## Introduction to Filters

- A class of Unix tools called *filters*.
  - Utilities that read from standard input, transform the file, and write to standard out
- Using filters can be thought of as *data oriented programming*.
  - Each step of the computation transforms data *stream*.

```
filter < abc > xyz
```



## Examples of Filters

- **Sort**
  - Input: lines from a file
  - Output: lines from the file sorted
- **Grep**
  - Input: lines from a file
  - Output: lines that match the argument
- **Awk**
  - Programmable filter

## cat: The simplest filter

- The `cat` command copies its input to output unchanged (*identity filter*). When supplied a list of file names, it **concatenates** them onto stdout.
- Some options:
  - `-n` number output lines (starting from 1)
  - `-v` display control-characters in visible form (e.g. `^C`)

---

```
cat file*
```

```
ls | cat -n
```

## head

- Display the first few lines of a specified file
- Syntax: `head [-n] [filename...]`
  - `-n` - number of lines to display, default is 10
  - `filename...` - list of filenames to display
- When more than one filename is specified, the start of each files listing displays  
`==>filename<==`


## tail

- Displays the last part of a file
- Syntax: `tail +|-number [lbc] [f] [filename]`  
or: `tail +|-number [l] [rf] [filename]`
  - `+number` - begins copying at distance `number` from beginning of file, if `number` isn't given, defaults to 10
  - `-number` - begins from end of file
  - `l,b,c` - `number` is in units of lines/block/characters
  - `r` - print in reverse order (lines only)
  - `f` - if input is not a pipe, do not terminate after end of file has been copied but loop. This is useful to monitor a file being written by another process

## head and tail examples

```
head /etc/passwd
head *.c
tail +20 /etc/passwd
ls -lt | tail -3
head -100 /etc/passwd | tail -5
tail -f /usr/local/httpd/access_log
```

## tee



- Copy standard input to standard output and one or more files
  - Captures intermediate results from a filter in the pipeline

## tee con't

- Syntax: `tee [-ai] file-list`
  - `-a` - append to output file rather than overwrite, default is to overwrite (replace) the output file
  - `-i` - ignore interrupts
  - `file-list` - one or more file names for capturing output
- Examples

```
ls | head -10 | tee first_10 | tail -5
who | tee user_list | wc
```

## Unix Text Files: Delimited Data

Tab Separated

Pipe-separated

John	99	COMP1011 2252424 Abbot, Andrew John  3727 1 M
Anne	75	COMP2011 2211222 Abdurjh, Saeed  3640 2 M
Andrew	50	COMP1011 2250631 Accent, Aac-Ek-Murhg  3640 1 M
Tim	95	COMP1021 2250127 Addison, Blair  3971 1 F
Arun	33	COMP4012 2190705 Allen, David Peter  3645 4 M
Sowmya	76	COMP4910 2190705 Allen, David Pater  3645 4 M

Colon-separated

```
root:ZHo1HAHZw8As2:0:0:root:/root:/bin/ksh
jas:nJz3ru5a/44Ko:100:100:John Shepherd:/home/jas:/bin/ksh
cs1021:iZ3s09005eZY6:101:101:COMP1021:/home/cs1021:/bin/bash
cs2041:rX9KwSSPqkLyA:102:102:COMP2041:/home/cs2041:/bin/csh
cs3311:mLRiCivmtI902:103:103:COMP3311:/home/cs3311:/bin/sh
```

## cut: select columns

- The **cut** command prints selected parts of input lines.
  - can select columns (assumes tab-separated input)
  - can select a range of character positions
- Some options:
  - **-f listOfCols**: print only the specified columns (tab-separated) on output
  - **-c listOfPos**: print only chars in the specified positions
  - **-d c**: use character *c* as the column separator
- Lists are specified as ranges (e.g. 1-5) or comma-separated (e.g. 2,4,5).


## cut examples

```
cut -f 1 < data
cut -f 1-3 < data
cut -f 1,4 < data
cut -f 4- < data
cut -d'|' -f 1-3 < data
cut -c 1-4 < data
```

*Unfortunately, there's no way to refer to "last column" without counting the columns.*

## paste: join columns

- The **paste** command displays several text files "in parallel" on output.
- If the inputs are files **a, b, c**
  - the first line of output is composed of the first lines of **a, b, c**
  - the second line of output is composed of the second lines of **a, b, c**
- Lines from each file are separated by a tab character.
- If files are different lengths, output has all lines from longest file, with empty strings for missing lines.



## paste example

```
cut -f 1 < data > data1
cut -f 2 < data > data2
cut -f 3 < data > data3
paste data1 data3 data2 > newdata
```

## sort: Sort lines of a file

- The **sort** command copies input to output but ensures that the output is arranged in ascending order of lines.
  - By default, sorting is based on ASCII comparisons of the whole line.
- Other features of **sort**:
  - understands text data that occurs in columns. (can also sort on a column other than the first)
  - can distinguish numbers and sort appropriately
  - can sort files "in place" as well as behaving like a filter
  - capable of sorting *very large* files

## sort: Options

- Syntax: **sort [-dfnr] [-o filename] [filename(s)]**
  - **-d** Dictionary order, only letters, digits, and whitespace are significant in determining sort order
  - **-f** Ignore case (fold into lower case)
  - **-t** Specify delimiter
  - **-n** Numeric order, sort by arithmetic value instead of first digit
  - **-r** Sort in reverse order
  - **-o filename** - write output to filename, filename can be the same as one of the input files
- Lots of more options...



## sort: Specifying fields

- Delimiter : `-t d`
- Old way:
  - `+f[.c] [options] [-f[.c] [options]`
    - `+2.1 -3 +0 -2 +3n`
  - Exclusive
  - Start from 0 (unlike cut, which starts at 1)
- New way:
  - `-k f[.c] [options] [,f[.c] [options]]`
    - `-k2.1 -k0,1 -k3n`
  - Inclusive
  - Start from 1

## sort Examples

```
sort +2nr < data
sort -k2nr data
sort -t: +4 /etc/passwd
sort -o mydata mydata
```

## uniq: list UNIQUE items

- Remove or report adjacent duplicate lines
- Syntax: `uniq [-cdu] [input-file] [output-file]`
  - `-c` Supersede the `-u` and `-d` options and generate an output report with each line preceded by an occurrence count
  - `-d` Write only the duplicated lines
  - `-u` Write only those lines which are not duplicated
  - The default output is the union (combination) of `-d` and `-u`

## wc: Counting results

- The word count utility, **wc**, counts the number of lines, characters or words
- Options:
  - `-l` Count lines
  - `-w` Count words
  - `-c` Count characters
- Default: count lines, words and chars

## wc and uniq Examples

```
who | sort | uniq -d
wc my_essay
who | wc
sort file | uniq | wc -l
sort file | uniq -d | wc -l
sort file | uniq -u | wc -l
```

## tr: Translate Characters

- Copies standard input to standard output with substitution or deletion of selected characters
- Syntax: `tr [-cds] [string1] [string2]`
  - `-d` delete all input characters contained in *string1*
  - `-c` complements the characters in *string1* with respect to the entire ASCII character set
  - `-s` squeeze all strings of repeated output characters in the last operand to single characters

## tr (continued)


- *tr* reads from standard input.
  - Any character that does not match a character in *string1* is passed to *standard output* unchanged
  - Any character that does match a character in *string1* is translated into the corresponding character in *string2* and then passed to *standard output*
- Examples
  - *tr s z* replaces all instances of *s* with *z*
  - *tr so zx* replaces all instances of *s* with *z* and *o* with *x*
  - *tr a-z A-Z* replaces all lower case characters with upper case characters
  - *tr -d a-c* deletes all a-c characters

## tr uses

- Change delimiter  
`tr '|' ':'`
- Rewrite numbers  
`tr ,. .,`
- Import DOS files  
`tr -d '\r' < dos_file`
- Find printable ASCII in a binary file  
`tr -cd '\na-zA-Z0-9' < binary_file`

## xargs

- Unix limits the size of arguments and environment that can be passed down to child
- What happens when we have a list of 10,000 files to send to a command?
- **xargs** solves this problem
  - Reads arguments as standard input
  - Sends them to commands that take file lists
  - May invoke program several times depending on size of arguments



## find utility and xargs

- `find . -type f -print | xargs wc -l`
  - `-type f` for files
  - `-print` to print them out
  - `xargs` invokes `wc l` or more times
- `wc -l a b c d e f g`  
`wc -l h i j k l m n o`  
...
- Compare to: `find . -type f -exec wc -l {} \;`

## Next Time

- Regular Expressions
  - Allow you to search for text in files
  - **grep** command
- We will soon learn how to write *scripts* that use these utilities in interesting ways.