

Kthreads, Mutexes, and Debugging

Sarah Diesburg
CS 3430
Operating Systems

Story of Kernel Development

Some context...

In the old days...

- There were no modules or virtual machines
- The kernel is a program
 - Has code, can compile, re-compile, make executable
 - When changes needed to be made, developers make changes in source and re-compile

How is the kernel different from a regular program?

- Mostly in how it is executed
 - Boot loader loads the kernel image/executable during boot time
 - Sets kernel mode
 - Jumps to the entry point in the image/executable
- Remember the generic booting sequence?

Quick Question

- How would you make changes to the kernel and run those changes?
 1. Make changes to the source
 2. Re-compile the kernel source
 3. Re-install the kernel source
 4. Make sure the bootloader sees the new kernel image (grub)
 5. Reboot and profit!

Getting more modern..

- Modules were created as bits of code that can be loaded and unloaded by the kernel in kernel mode
- Made development easier
 - Instead of re-compiling, re-installing, and rebooting into the new kernel, one could just re-compile and load a module

Quick Question

- How would you make changes to a module and run those changes?
 1. Make changes to module source code
 2. Re-compile the module
 3. Load the new module

Present Day

- Reboots into new kernels and loading new modules often freezes machines
- Enter virtual machine software
 - **Process** that emulates the hardware necessary to run an OS in user-space
 - Guest OS is executed inside the virtual machine process!

Kthreads

Run the main logic of your module in a kthread!

Refresher: hello.c

```
#include <linux/init.h>
#include <linux/module.h>
MODULE_LICENSE("Dual BSD/GPL");

static int hello_init(void)
{
    printk(KERN_ALERT "Hello, world!\n");
    return 0;
}

static void hello_exit(void)
{
    printk(KERN_ALERT "Goodbye, sleepy world.\n");
}

module_init(hello_init);
module_exit(hello_exit);
```

Kernel Modules

- Remember, kernel modules are very ***event-based***
- We need a way to start an independent thread of execution in response to an event
 - e.g. `start_kitchen()` for project 2...

kthread_run

```
kthread_run(threadfn, data, namefmt, ...)
```

- Creates a new thread and tells it to run
 - `threadfn` – the name of the function the thread should run
 - `data` – data pointer for `threadfn` (can be NULL if the function does not take any args)
 - `namefmt` – name of the thread (displayed during “ps” command)
- Returns a `task_struct`

kthread_run example

```
struct task_struct *t;  
  
t = kthread_run(run, NULL, "my_elevator");  
    if (IS_ERR(t)){  
        ret=PTR_ERR(t);  
    }
```

kthread_stop

```
int kthread_stop(struct task_struct * k);
```

- Sets `kthread_should_stop` for `k` to return true, wakes the thread, and waits for the thread to exit
- Returns the result of the thread function

kthread_stop_example

```
ret=kthread_stop(t);  
if(ret != -EINTR)  
    printk("Main logic tread stopped.\n");
```

Thread Function Example

```
static int run(void *arg)
{
    /* Lock here */

    while(!kthread_should_stop()) {

        /* Do stuff */

    }

    /* Unlock here */
    printk("%s: kernel thread exits.\n",
__FUNCTION__);
    return 0;
}
```


Sample kthread code

- Take a look at the sample module that spawns a kthread on load
 - You will want to move this kthread to start when user writes a “0” to /proc/kitchen
 - You will want to stop the kthread when user writes a “-1” to /proc/kitchen

Concurrency Aspects of Project 2

- Synchronizing access to request queue(s)
 - Multiple producers may access request queue(s) at the same time
 - Multiple consumers may access request queue(s) at the same time
 - Synchronizing access to other global data
-

Kitchen Queue Concurrency

- Orders may appear on the queue at the same time the kitchen module checks the queue
- The status may be read at the same time that you're updating
 - Number of orders that you've serviced
 - Which order is currently being processed
 - Which slot in the queue kitchen is looking at
- How do you guarantee correctness?

Global Data vs. Local Data

- ***Global data*** is declared at global scope, e.g. outside of any function body
 - Often necessary for kernel programming
- Particularly sensitive to concurrency issues
 - Be **extra** careful when handling globals

Global Data vs. Local Data

- ***Local data*** is declared within a function
- Local data is sensitive to concurrency issues when it depends on global data or when parallel access is possible
 - Think carefully about whether it needs to be synchronized

Synchronization Primitives

- Semaphores
 - User space
 - Kernel space
- Mutexes
 - User space
 - Kernel space
- Spin locks
 - Kernel space
- Atomic Functions

Synchronization Primitives

(We'll Only Cover These)

- Mutexes
 - User space
 - Kernel space
- Does anyone remember the differences between a mutex and semaphore?

The Mutex

Caught up in the mutex?

Mutexes

- Mutex – A construct to provide MUTual EXclusion
- Based on the concept of a semaphore
- Found at
`<source_dir>/include/linux/mutex.h`
- Can be locked or unlocked
 - Only one thread of execution may hold the lock at a time

Kernel-Space Mutex - Initialization

```
■ mutex_init(&mutex)
```

- Declare and initialize a mutex
 - Only initialize it once

Kernel-Space Mutex - Locking

```
■ void mutex_lock(struct mutex *);  
■ int  mutex_lock_interruptible(struct mutex *);
```

- mutex_lock() can wait indefinitely
- mutex_lock_interruptible() locks a mutex as long as it is not interrupted
 - returns 0 if locking succeeded, < 0 if interrupted
- Use of the interruptible version is typically preferred

Kernel-Space Mutex – Unlocking

```
■ void mutex_unlock(struct mutex *);
```

- Guarantees that the mutex is unlocked
 - Why is there no interruptible version of this function?

Kernel-Space Mutex Example

```
/* Declare your mutex */
struct mutex my_mutex;

/* Initialize your mutex */
mutex_init(&my_mutex);

/* Lock */
if(mutex_lock_interruptible(&my_mutex))
    return -ERESTARTSYS;

/* Do stuff to protected global variables */

/* Unlock */
mutex_unlock(&my_mutex);
```

User-Space Mutex

- Also used with pthreads in regular user applications
 - pthreads operate very similar to kthreads
 - Might be useful if you are prototyping your elevator in user-space before porting to kernel

User-Space Mutex - Initialization

```
■ int pthread_mutex_init(pthread_mutex_t *, NULL);  
■ int pthread_mutex_destroy(pthread_mutex_t *);
```

- Pthread_mutex_init() dynamically allocates a mutex
- Pthread_mutex_destroy() frees a mutex

User-Space Mutex - Locking

```
■ int pthread_mutex_lock(pthread_mutex_t *);
```

- Returns 0 on locking, < 0 otherwise

User-Space Mutex - Unlocking

```
■ int pthread_mutex_unlock(pthread_mutex_t *);
```

- Returns 0 on unlocking, < 0 otherwise

Kitchen Scheduling Advice

General Advice

- Just make kitchen work first
 - Use a very simple algorithm
 - My kitchen search the queue in round-robin fashion and processes every Beef Wellington
- Optimize if there is time

Round Robin

- Method:
 - Service requests in round-robin fashion (e.g. queue slot 0, 1, 2, 3, etc.)
- Pros/Cons?

Shortest Job First

- Method:
 - Service fastest orders first
- Pros/Cons?

SCAN

- Method:
 - Service requests in one direction, then go backwards
- Pros/Cons?

Hybrid

- Combine methods, come up with something new
- Up to your creativity

Project Demos and Deliverables

Basic Information

- Project 2 is due on October 26th
 - Due at demo time
- Please sign up for a time to demo your elevator project
- If you wish to use slack days, make an appointment with me on the day you wish to turn in your project

Project Deliverables

- Before final demo, please zip and upload the following to eLearning
 - README
 - Part 2 – source and Makefile
 - Part 3 – source and Makefile
- You may have to copy the files to vermin or your workstation to zip them (or else read the man page on 'tar' to zip on your virtual machine).

Demo

- Will only have time to demo Part 3 – elevator
- Please look at grading sheet to understand what I will be looking for

Getting Help

- Sign up for a halfway demo
- Regular office hours

Other Hints

- This is not a simple project
 - ❑ Setup is different
 - ❑ You need to use different files and methods of compilation/running
 - ❑ Do NOT wait until 2 days before it is due to start
 - Too late
 - ❑ The Internet will likely NOT be very helpful

Other Hints

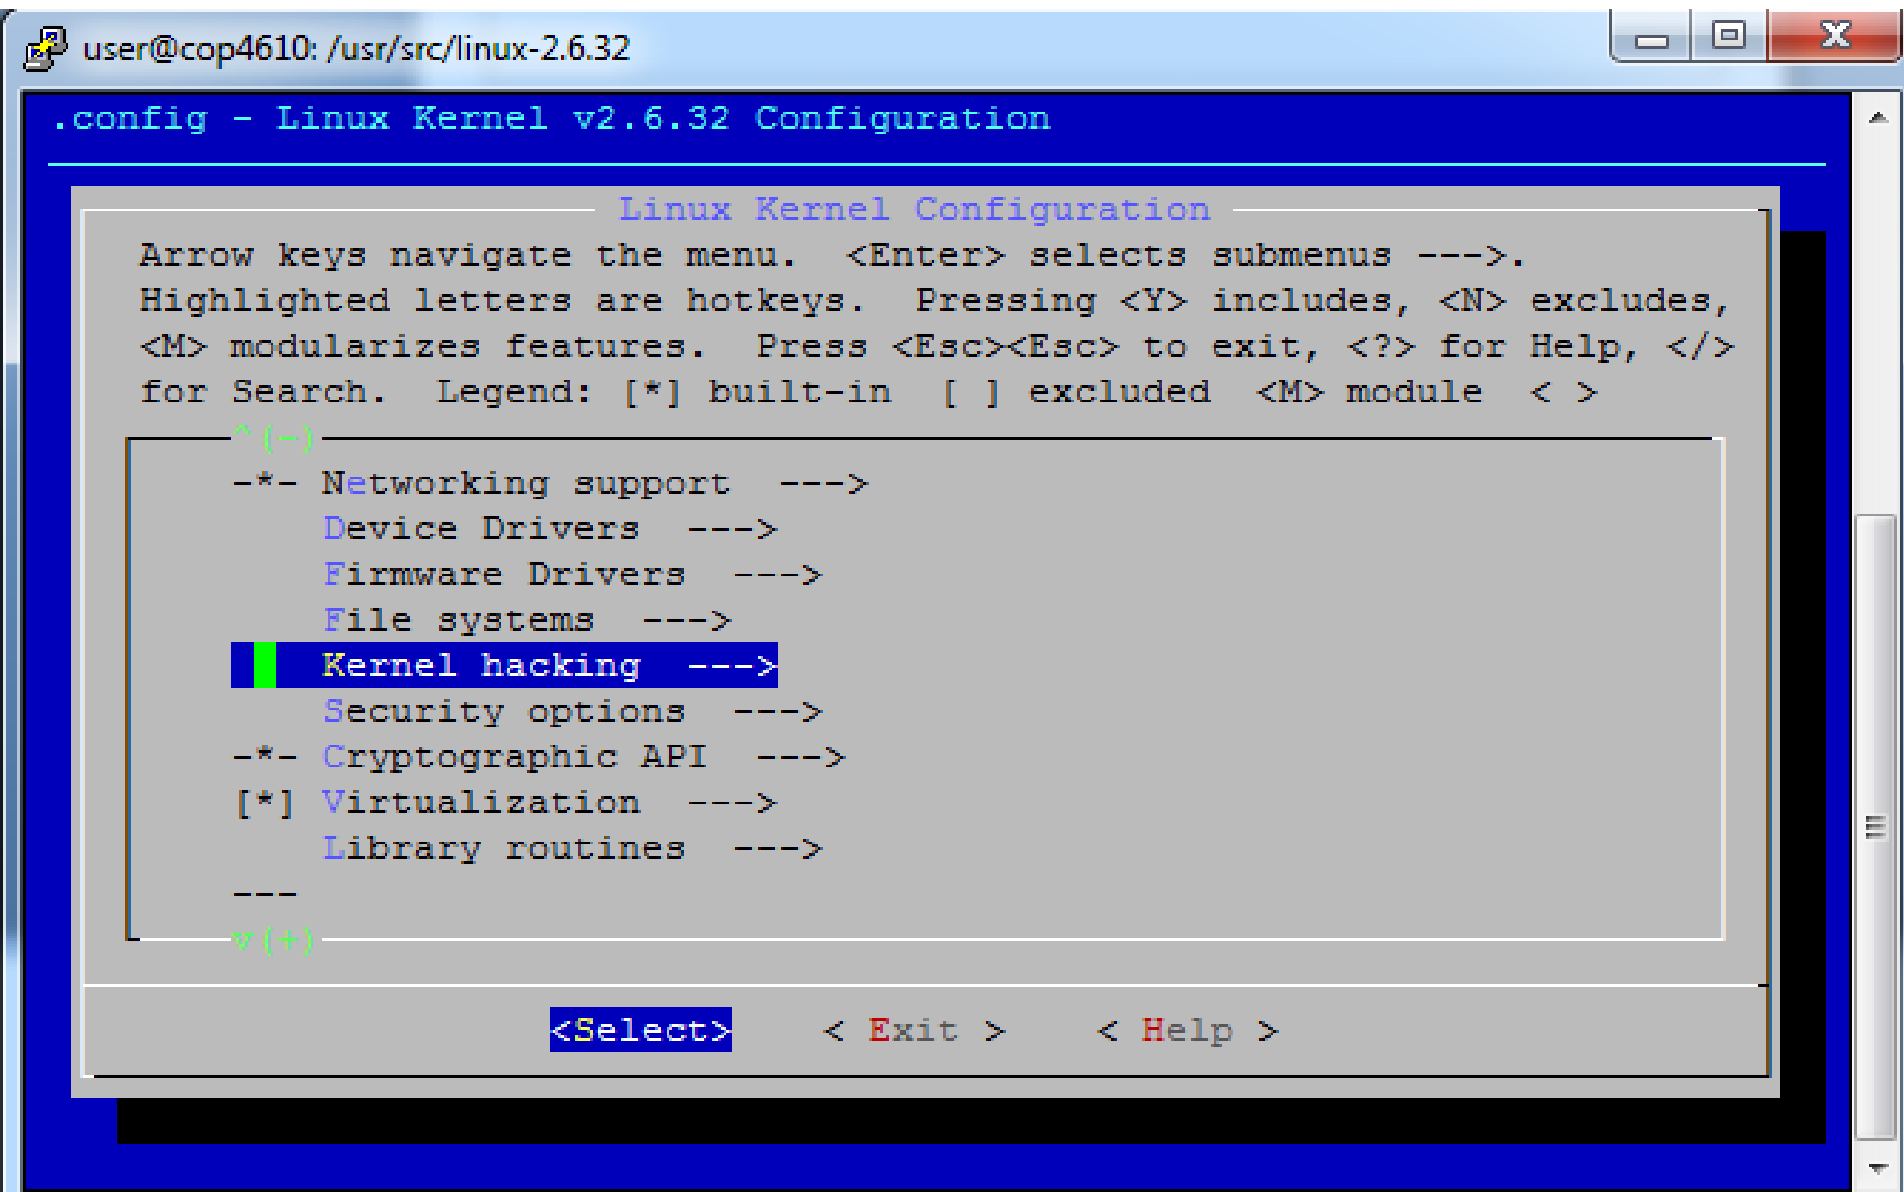
- Set milestones over the next few weeks
- Ask questions early
 - If it's a good question, I'll share it with the class

Debugging

Kernel Debugging Configurations

- Timing info on printks
- __deprecated logic
- Detection of hung tasks
- SLUB debugging
- Kernel memory leak detector
- Mutex/lock debugging
- Kmemcheck
- Check for stack overflow
- Linked list debugging

Select Kernel Hacking



```
user@cop4610: /usr/src/linux-2.6.32

.config - Linux Kernel v2.6.32 Configuration

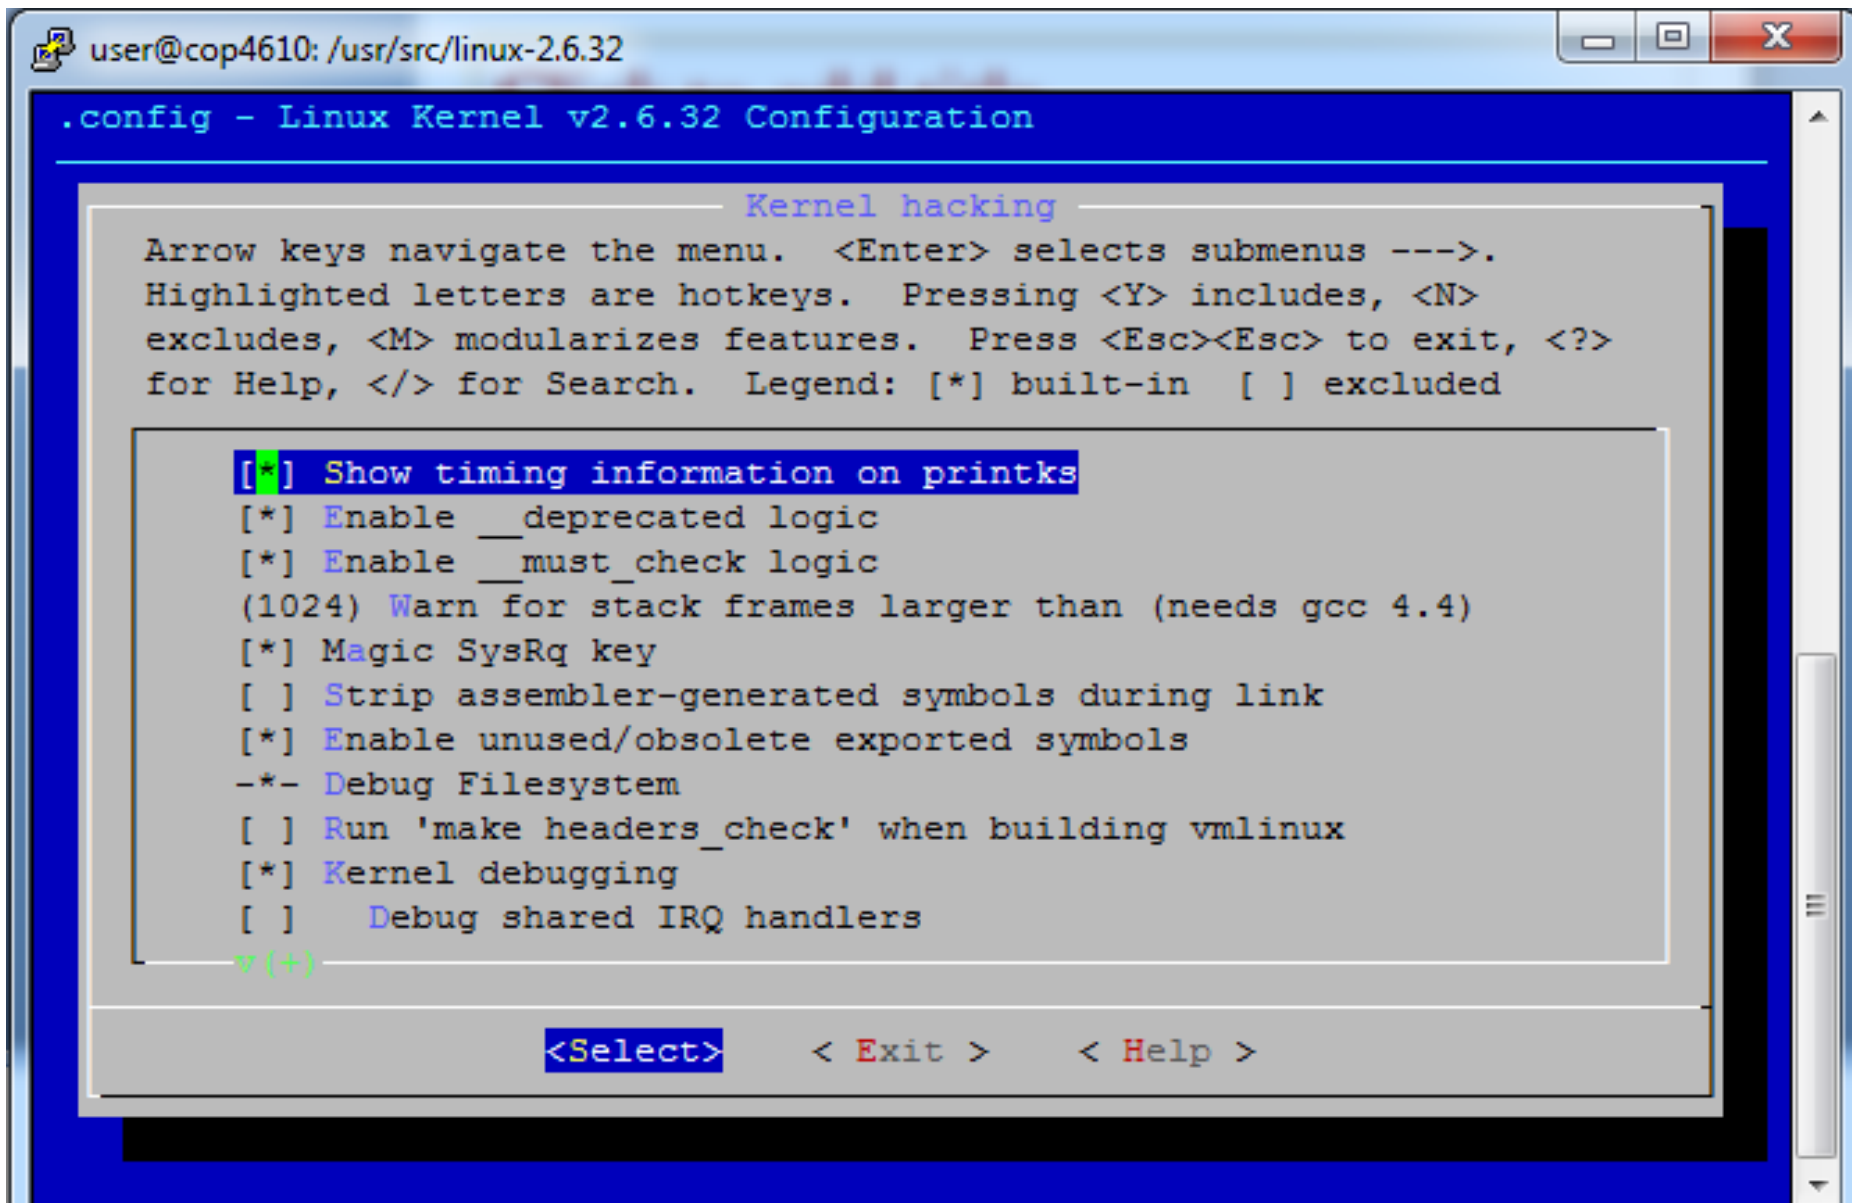
Linux Kernel Configuration

Arrow keys navigate the menu.  <Enter> selects submenus --->.
Highlighted letters are hotkeys.  Pressing <Y> includes, <N> excludes,
<M> modularizes features.  Press <Esc><Esc> to exit, <?> for Help, </>
for Search.  Legend: [*] built-in  [ ] excluded  <M> module  < >

^(-)
-- Networking support --->
  Device Drivers --->
  Firmware Drivers --->
  File systems --->
  Kernel hacking --->
  Security options --->
-- Cryptographic API --->
[*] Virtualization --->
  Library routines --->
---
v(+)
```

<Select> <Exit> <Help>

Enable Debugging Options



The screenshot shows a terminal window titled "user@cop4610: /usr/src/linux-2.6.32". The main title of the menu is ".config - Linux Kernel v2.6.32 Configuration". The current section is "Kernel hacking". A paragraph of instructions explains the navigation: "Arrow keys navigate the menu. <Enter> selects submenus --->. Highlighted letters are hotkeys. Pressing <Y> includes, <N> excludes, <M> modularizes features. Press <Esc><Esc> to exit, <?> for Help, </> for Search. Legend: [*] built-in [] excluded". A list of options is shown, with "[*] Show timing information on printk" highlighted in blue. Other options include enabling deprecated logic, must-check logic, warning for stack frames, Magic SysRq key, stripping assembler-generated symbols, enabling unused/obsolete exported symbols, debugging the filesystem, running 'make headers_check', kernel debugging, and debug shared IRQ handlers. At the bottom, navigation keys are listed: "<Select>", "< Exit >", and "< Help >".

```
user@cop4610: /usr/src/linux-2.6.32

.config - Linux Kernel v2.6.32 Configuration

                                Kernel hacking

Arrow keys navigate the menu.  <Enter> selects submenus --->.
Highlighted letters are hotkeys.  Pressing <Y> includes, <N>
excludes, <M> modularizes features.  Press <Esc><Esc> to exit, <?>
for Help, </> for Search.  Legend: [*] built-in  [ ] excluded

[*] Show timing information on printk
[*] Enable __deprecated logic
[*] Enable __must_check logic
(1024) Warn for stack frames larger than (needs gcc 4.4)
[*] Magic SysRq key
[ ] Strip assembler-generated symbols during link
[*] Enable unused/obsolete exported symbols
-- Debug Filesystem
[ ] Run 'make headers_check' when building vmlinux
[*] Kernel debugging
[ ]   Debug shared IRQ handlers
v(+)-

<Select>    < Exit >    < Help >
```

Debugging through reads to `/proc/kitchen`

- Necessary to help you “see” what’s going on!
- General process
 - Identify data to monitor in your module
 - Run your module
 - Query `/proc/kitchen` for that information at any time

Kernel Oops and Other Errors

- Kernel errors often only appear on first tty (terminal interface)
 - Why?

```
done.
Unmounting local filesystems...umount: tmpfs busy - remounted read-only
done.
Unable to handle paging request at virtual address f8fb37dc
printing error
c02bbf60
Oops: 3774f067
Oops: 0000 [#1]
PREEMPT
Modules linked in: bnep rfcomm hidp l2cap irda crc_ccitt binfmt_misc ipu6 fire
ec snd_pcm_oss snd_mixer_oss snd_pcm snd_timer snd soundcore snd_page_alloc
hci_hcd uhci_hcd tg3 ohci1394 yenta_socket rsrc_nonstatic pcmcia_core nls_iso
ooth wbsd mmc_block mmc_core tun msr cpuid cpufreq_stats container video hotk
rnal battery ac speedstep_centrino freq_table processor sr_mod sbp2 scsi_mod
CPU: 0
EIP: 0060:[<c02bbf60>] Tainted: P ULI
EFLAGS: 00010282 (2.6.13-rc5-x300)
EIP is at suspend_device+0xa8/0x17b
eax: f8fb3640 ebx: f71b8be4 ecx: 00000000 edx: 00000000
esi: f71b8be4 edi: 00000000 ebp: 00000003 esp: f680de44
ds: 007b es: 007b ss: 0068
Process halt (pid: 7171, threadinfo=f680c000 task=f6d08020)
Stack: c038c83f 00000066 f680de6c c011d134 c1a06aa0 00000246 f71b8ca4 f71b8ca
f71b8d3c f71b8be4 00000000 00000003 c02bc102 f71b8be4 00000003 0000000
4321fedc bf8b9f29 b7f88e80 f680c000 c0139e94 00000003 00000003 0000000
Call Trace:
[<c011d134>] activate_task+0x61/0x70
[<c02bc102>] device_suspend+0xcf/0x1d9
[<c0139e94>] kernel_power_off+0x35/0x4e
```

```
done.
Unmounting local filesystems...umount: tmpfs busy - remounted read-only
done
Unable to handle kernel paging request at virtual address f8fb37dc
printing eip:
c02bbf60
*pde = 3774f067
Oops: 0000 [#1]
PREEMPT
Modules linked in: bnep rfcomm hidp l2cap irda crc_ccitt bluetooth_misc ipu6 fire
ec snd_pcm_oss snd_mixer_oss snd_pcm snd_timer snd soundcore snd_page_alloc f
hci_hcd uhci_hcd tg3 ohci1394 yenta_socket rsrc_nonstatic s_is
ooth wbsd mmc_block mmc_core tun nsr cpuid cpufreq_st
rml battery ac speedstep_centrino freq_table processo
CPU: 0
EIP: 0060:[<c02bbf60>] Tainted: P ULI
EFLAGS: 00010282 (2.6.13-rc5-x300)
EIP is at suspend_device+0xa8/0x17b
eax: f8fb3640 ebx: f71b8be4 ecx: 00000000 edx: 00000000
esi: f71b8be4 edi: 00000000 ebp: 00000003 esp: f680de44
ds: 007b es: 007b ss: 0068
Process halt (pid: 7171, threadinfo=f680c000 task=f6d08020)
Stack: c038c83f 00000066 f680de6c c011d134 c1a06aa0 00000246 f71b8ca4 f71b8ca
f71b8d3c f71b8be4 00000000 00000003 c02bc102 f71b8be4 00000003 0000000
4321fedc bf8b9f29 b7f88e80 f680c000 c0139e94 00000003 00000003 0000000
Call Trace:
[<c011d134>] activate_task+0x61/0x70
[<c02bc102>] device_suspend+0xcf/0x1d9
[<c0139e94>] kernel_power_off+0x35/0x4e
```

Reason
for failure

```
done.
Unmounting local filesystems...umount: tmpfs busy - remounted read-only
done.
Unable to handle kernel paging request at u...b37dc
printing eip:
c02bbf60
*pde = 3774f067
Oops: 0000 [#1]
PREEMPT
Modules linked in: bnep rfcomm hidp l2cap irda crc_ccitt binfmt_misc ipu6 fir
ce snd_pcm oss snd_mixer_oss snd_pcm snd_timer snd soundcore snd_page_alloc
hci_hcd uhci_hcd tg3 ohci1394 yenta_socket rsrc_nonstatic pcmcia_core nls_iso
ooth wbsd mmc_block mmc_core tun msr cpuid cpufreq_stats container video hotk
rnal battery ac speedstep_centrino freq_table processor sr_mod sbp2 scsi_mod
CPU: 0
EIP: 0060:[<c02bbf60>] Tainted: P ULI
EFLAGS: 00010282 (2.6.13-rc5-x300)
EIP is at suspend_device+0xa8/0x17b
eax: f8fb3640 ebx: f71b8be4 ecx: 00000000 edx: 00000000
esi: f71b8be4 edi: 00000000 ebp: 00000003 esp: f680de44
ds: 007b es: 007b ss: 0068
Process halt (pid: 7171, threadinfo=f680c000 task=f6d08020)
Stack: c038c83f 00000066 f680de6c c011d134 c1a06aa0 00000246 f71b8ca4 f71b8ca
f71b8d3c f71b8be4 00000000 00000003 c02bc102 f71b8be4 00000003 0000000
4321fedc bf8b9f29 b7f88e80 f680c000 c0139e94 00000003 00000003 0000000
Call Trace:
[<c011d134>] activate_task+0x61/0x70
[<c02bc102>] device_suspend+0xcf/0x1d9
[<c0139e94>] kernel_power_off+0x35/0x4e
```

Current
drivers


```
done.
Unmounting local filesystems...umount: tmpfs busy - remounted read-only
done.
Unable to handle kernel paging request at virtual address f8fb37dc
printing eip:
c02bbf60
*pde = 3774f067
Oops: 0000 [#1]
PREEMPT
Modules linked in: bnep rfcomm hidp l2cap irda crc_ccitt binfmt_misc ipv6 fire
ec snd_pcm_oss snd_mixer_oss snd_pcm snd_timer snd soundcore snd_page_alloc f
hci_hcd uhci_hcd tg3 ohci1394 yenta_socket rsrc_nonstatic pcmcia_core nls_iso
ooth wbsd mmc_block mmc_core tun msr cpuid cpufreq_stats container video hotk
rnal battery ac speedstep_centrino freq_table processor sr_mod sbp2 scsi_mod
CPU: 0
EIP: 0060:[<c02bbf60>] Tainted: P ULI
EFLAGS: 00010282 (2.6.13-rc5-x300)
EIP is at suspend_device+0xa8/0x17b
eax: f8fb3640 ebx: f71b8be4 ecx: 00000000
esi: f71b8be4 edi: 00000000 ebp: 00000000
ds: 007b es: 007b ss: 0068
Process halt (pid: 7171, thread: f680c000 task=f6d08020)
Stack: c038c83f 00000066 f680de6c c011d134 c1a06aa0 00000246 f71b8ca4 f71b8ca
f71b8d3e f71b8be4 00000000 00000003 c02bc102 f71b8be4 00000003 00000000
4321fedc bf8b9f29 b7f88e80 f680c000 c0139e94 00000003 00000003 00000000
Call Trace:
[<c011d134>] activate_task+0x61/0x70
[<c02bc102>] device_suspend+0xcf/0x1d9
[<c0139e94>] kernel_power_off+0x35/0x4e
```

Call Trace


```
hsd mmc_block mmc_core tun msr cpuid cpufreq_stats container video hotkey fan
battery ac speedstep_centrino freq_table processor sr_mod sbp2 scsi_mod ieee13
0
0060:[<c02bbf60>]    Tainted: P          ULI
: 00010282    (2.6.13-rc5-x300)
: at suspend_device+0xa8/0x17b
8fb3640    ebx: f71b8be4    ecx: 00000000    edx: 00000000
71b8be4    edi: 00000000    ebp: 00000003    esp: f680de44
7b    es: 007b    ss: 0068
s halt (pid: 7171, threadinfo=f680c000 task=f6d08020)
c038c83f 00000000 f680de6c c011d134 c1a06aa0 00000246 f71b8ca4 f71b8cec
f71b8a3c f71b8be4 00000000 00000003 c02bc102 f71b8be4 00000003 00000003
4321fedc bf8b9f29 b7f88e80 f680c000 c0139e94 00000003 00000003 00000000
Trace:
1d134>] activate_task+0x61/0x70
bc102>] device_suspend+0xcf/0x1d9
39e94>] kernel_power_off+0x35/0x4e
3a02e>] sys_reboot+0x181/0x1af
3465f>] __group_send_sig_info+0xcb/0xe9
78c25>] preempt_schedule+0x4a/0x56
34a9b>] kill_proc_info+0x69/0x6b
38335>] sys_kill+0x5b/0x62
a3bcd>] do_ioctl+0x2d/0x81
a3db0>] ufs_ioctl+0x61/0x1fb
a3f86>] sys_ioctl+0x3c/0x5a
03c35>] syscall_call+0x7/0xb
85 54 01 00 00 85 c0 74 0e 8b 88 44 01 00 00 85 c9 0f 85 84 00 00 00 89 93 48
0 <8b> 80 8c 01 00 00 85 c0 0f 84 c1 00 00 00 85 d2 0f 85 b9 00 00
rc0.d/S90halt: line 48: 7171 Segmentation fault      halt -d -f -i $poweroff
```

Call Trace

```
0000: (c02b8f60)ainted: F 0L1
: 00010282 (2.6.13-rc5-x300)
at suspend_device+0xa8/0x17b
8fb3640 ebx: f71b8be4 ecx: 00000000 edx: 00000000
71b8be4 edi: 00000000 ebp: 00000003 esp: f680de44
7b es: 007b ss: 0068
s halt (pid: 7171, threadinfo=f680c000 task=f6d08020)
c038c83f 00000066 f680de6c c011d134 c1a06aa0 00000246 f71b8ca4 f71b8cec
f71b8d3c f71b8be4 00000000 00000003 c02bc102 f71b8be4 00000003 00000003
4321fedc bf8b9f29 b7f88e80 f680c000 c0139e94 00000003 00000003 00000000
race:
1d134>] activate_task+0x61/0x70
bc102>] device_suspend+0xcf/0x1d9
39e94>] kernel_power_off+0x35/0x4e
3a02e>] sys_reboot+0x181/0x1af
3465f>] __group_send_sig_info+0xcb/0xe9
78c25>] preempt_schedule+0x4a/0x56
34a9b>] kill_proc_info+0x69/0x6b
38335>] sys_kill+0x5b/0x62
a3bcd>] do_ioctl+0x2d/0x81
a3db0>] ufs_ioctl+0x61/0x1fb
a3f86>] sys_ioctl+0x3c/0x5a
03c35>] syscall_call+0x7/0xb
83 54 01 00 00 85 c0 74 0e 8b 88 44 01 00 00 85 c9 0f 85 84 00 00 00 89 93 48
(0b) 80 9c 01 00 00 85 c0 0f 84 c1 00 00 00 85 d2 0f 85 b9 00 00
rc0.d/S90halt: line 48: 7171 Segmentation fault halt -d -f -i $poweroff
SS: Unregistering class device, ID = 'ucs1'
hotplug - name = ucs1
e class 'ucs1': release.
Unregistering class device, ID = 'ucs1'
```



**Failed
command**