

Debugging Multi-threaded Programs

Designing for debugging

- Use proper software development practices
 - Don't code first without designing first
 - Design with testing in mind
- More crucial for multi-threaded applications
 - multi-threaded applications are inherently more complicated
 - Large number of special cases
 - Wide range of possible paths
 - Access patterns may vary widely between single and multi-core
 - Multi-threaded bugs may not appear when running under debugger
 - Bugs sensitive to timing of events

Points to keep in mind

- Design the application so it can run sequentially
 - Validate in this mode first
- Use established parallel programming patterns
 - Solve many common problems
- Include built-in debug support in the application
 - Useful to be able to examine state of system at arbitrary time
 - Add functions to display state of thread or all threads
 - Use trace buffers to record sequence of accesses to shared resource

Programming with OpenMP*

3

Do Code reviews

- Reviews challenging for parallel programs
- One technique
 - Individual reviewers examine from perspective of one of threads
 - Step through events as thread would
 - Have reviewers take and release shared resources
- Why validate?
 1. Take a few weeks up front to validate and verify the design
 2. Have to redesign from scratch when doesn't scale
- Why review?
 - Do you want to be debugging unpredictable bugs a week before the demonstration

Programming with OpenMP*

4

Trace Buffers for Debugging Synchronization Bugs

- Need 2 pieces of information
 1. Which threads are accessing resource at time of failure
 2. When access took place
- A log or trace of the access pattern of different threads helps narrow down code to be reviewed
 - Trace buffer one way to do this
 - Mechanism for logging events
 - Uses atomic counter to keep track of empty slot in event log array
 - Information stored up to developer

Programming with OpenMP*

5

Windows Trace Buffer example

- Windows app - progs\Ahkter\Ch8\DeadlockDebugApp.cpp
- Trace buffer stores 1024 events in circular buffer
 - Atomic index wraps
 - Don't need to dynamically resize buffer
- traceBufferElement is the event descriptor
- Three operations implemented
 - InitializeTraceBuffer
 - Initialized atomic counter to -1
 - AddEntryToTraceBuffer
 - PrintTraceBuffer
 - Useful if debugger allows execution of code at breakpoint
 - GDB and Visual Studio allow this
 - Can see recent events

Programming with OpenMP*

6

Issues with trace buffers

- Logs events as passed into buffer
 - Does not guarantee that will log events exactly as they occur in times
 - Can have data races between thread writing shared variables and trace buffer
- Example

Thread1	Thread2
m_global=dowork();	
	Thread_local_data=m_global;
	AddEntryToTraceBuffer(msg);
AddEntryToTraceBuffer(msg);	

- Buffer does not reflect sequence of events
- Can solve by placing locks around event and related log to trace buffer

Programming with OpenMP*

7

Drawbacks to protecting logging

- May mask synchronization problems in original code
- Have protected the critical section that was unprotected
- When tracking down race condition, should avoid synchronization to access buffer
 - Preferred method is to log before and after event
 - If before and after messages occur in order can assume the event ordered
 - If before and after messages are interleaved then the order of events is indeterminate
- In general if you protect/synchronize access and application works, probably a problem with synchronization in the original code

Programming with OpenMP*

8

Debugging using gdb

- Gdb capabilities for multi-threaded debugging
- Not all versions support all features
 - Automatic notification when threads created
 - displays threadid called **systag** - OS identification of thread
 - List all threads
 - **info threads** will print list of threads giving GDB thread number, systag, current stack frame, indicates active thread with *
 - thread-specific breakpoints
 - **break linespec thread threadnum**
 - Stops all threads
 - Can have conditional break-points e.g.
break buffer.c:33 thread 7 if level > watermark

Programming with OpenMP*

9

Debugging using gdb

- thread-specific breakpoints - other issues
 - System calls may return early when a breakpoint is triggered
 - GDB uses signals to manage breakpoints, cause system calls to return
 - Advisable to check return values of sys calls and handle this case
 - Ex: sleep(30) syscall may return early – return value will be number of seconds left to sleep
 - GDB does not singlestep all threads in lockstep
 - may execute lots of code in other threads
 - If breakpoints in other threads may jump to these
 - Some versions of GDB have scheduler locker – only current thread can run

Programming with OpenMP*

10

Debugging using gdb

- Ability to switch between threads
 - Use **thread threadnum** command
- Ability to apply commands to group of threads
 - Use command **apply**
 - argument can be threadnum or keyword **all**

Programming with OpenMP*

11