

OS Lab Tutorial

Thread

Today ...

- Pthreads
 - Create/Terminate
 - Obtain the returned value
- Synchronization
 - mutex
 - cond
- Some Programming Tips
- Assignment Spec

Threads VS. Processes

Platform	fork()			pthread_create()		
	real	user	sys	real	user	sys
Intel 2.6 GHz Xeon E5-2670 (16 cores/node)	8.1	0.1	2.9	0.9	0.2	0.3
Intel 2.8 GHz Xeon 5660 (12 cores/node)	4.4	0.4	4.3	0.7	0.2	0.5
AMD 2.3 GHz Opteron (16 cores/node)	12.5	1	12.5	1.2	0.2	1.3
AMD 2.4 GHz Opteron (8 cores/node)	17.6	2.2	15.7	1.4	0.3	1.3
IBM 4.0 GHz POWER6 (8 cpus/node)	9.5	0.6	8.8	1.6	0.1	0.4
IBM 1.9 GHz POWER5 p5-575 (8 cpus/node)	64.2	30.7	27.6	1.7	0.6	1.1
IBM 1.5 GHz POWER4 (8 cpus/node)	104.5	48.6	47.2	2.1	1	1.5
INTEL 2.4 GHz Xeon (2 cpus/node)	54.9	1.5	20.8	1.6	0.7	0.9
INTEL 1.4 GHz Itanium2 (4 cpus/node)	54.5	1.1	22.2	2	1.2	0.6

Timings reflect 50,000 process/thread creations, units in seconds

Data Source: <https://computing.llnl.gov/tutorials/pthreads/> Code: https://computing.llnl.gov/tutorials/pthreads/fork_vs_thread.txt

POSIX Threads (pthreads)

- POSIX - Portable Operating System Interface
 - maintain the compatibility between operating system
 - The standard, *POSIX.1c, Threads extensions (IEEE Std 1003.1c-1995)*, defines an **API** for creating and manipulating threads.
- Standard Unix threading library
- Provides useful concurrent constructs
 - mutex
 - conditional variables
- Note
 - The POSIX semaphore API works with POSIX threads but is not part of threads standard. The semaphore procedures are prefixed by "sem_" instead of "pthread_" (http://en.wikipedia.org/wiki/POSIX_Threads)

pthread routines

- `pthread_attr_init`, `pthread_attr_destroy`
 - Initialize/destroy the attribute object of a thread
 - joinable/detached
- `pthread_create`
 - Create a thread
- `pthread_join`
 - wait for threads to finish
- `pthread_exit`
 - finish a thread
- `pthread_cancel`
 - Cancel a thread

Found no Man Page Info. ?

- Man Pages

- The manual documents for POSIX pthread library may **not** be installed on some versions of the Linux system. You may have to use the online version or install it by yourself:

- Linux Man Pages: <http://linux.die.net/man/>

- Search by "Names" = pthread

- Installation (Ubuntu):

- `$ sudo apt-get install manpages-posix-dev`
- `$ sudo apt-get install glibc-doc`

- Then check the installation by:

- `$ man pthread_create`

pthread_create

- `<pthread.h>`
- `int pthread_create(pthread_t *restrict thread, const pthread_attr_t *restrict attr, void *(*start_routine)(void*), void *restrict arg);`
- Arguments
 - pointer to a thread ID, type: `pthread_t`
 - attribute object of a thread (could be **NULL**)
 - routine to be executed in a thread
 - arguments for the routine (could be **NULL**)

// Create a detached thread

`pthread_attr_t attr; // thread attribute`

// set thread detachstate attribute to DETACHED

`pthread_attr_init(&attr);`

`pthread_attr_setdetachstate(&attr, PTHREAD_CREATE_DETACHED);`

// create the thread

`pthread_create(&tid, &attr, start_routine, arg);`

pthread_create

```
// From apue 2
#include <stdio.h>
#include <stdlib.h>
#include <pthread.h>

void *print_sth( void *ptr ){
    char* message;
    message = (char*)ptr;
    printf("%s\n",message);
    return NULL;
}
```

```
int main()
{
    pthread_t thread1;
    char *message1 = "Thread 1";
    iret1 = pthread_create( &thread1, NULL, print_sth,
                           (void*) message1);

    pthread_join( thread1, NULL);
    printf("Thread 1 finished.\n");
    return 0;
}
```


pthread_join

- Wait for a (joinable) thread to finish. If return value is **not NULL**, it copies the value into a location pointed to by ***retval**
- `<pthread.h>`
- `int pthread_join(pthread_t thread, void **retval);`
- Arguments
 - id of the created thread
 - returned value from the created thread

pthread_join

```
// From apue 2
#include <stdio.h>
#include <stdlib.h>
#include <pthread.h>

void *print_sth( void *ptr ){
    char* message;
    message = (char*)ptr;
    printf("%s\n",message);
    return NULL;
}
```

```
int main()
{
    pthread_t thread1;
    char *message1 = "Thread 1";
    iret1 = pthread_create( &thread1, NULL, print_sth,
                           (void*) message1);

    // Blocked. Until thread1 finishes
    pthread_join( thread1, NULL);
    printf("Thread 1 finished.\n");
    return 0;
}
```

pthread_exit and pthread_cancel

- `<pthread.h>`
- `void pthread_exit(void *retval);`
 - Terminate the calling thread and return the value
- `int pthread_cancel(pthread_t thread)`
 - Send a cancellation request to a thread

Programming Notification

- Include the corresponding header files
 - `#include <pthread.h>`
 - And others, ... e.g. "pthread_t" (id of a thread) is defined in `<sys/types.h>`
- Compile
 - `$ gcc -Wall -o test pthread_test.c -lpthread`
- Check the return value of a function call
 - Many functions write/record/set error information in a global variable `errno` when returning from errors
 - You can view the error message by examining `errno`

Errno, perror, strerror

- **errno** is an **int** variable declared in **<errno.h>**
- Each valid integer number (**nonzero**) represents an error code.
 - (\$ man errno)

2	ENOENT	No such file or directory
---	--------	---------------------------

- A good programming style is to ***always*** check the return value of a function and the associated **errno**.
- `perror()` / `strerror()` can be used to translate the number stored in `errno` into a human-readable string, i.e., *No such file or directory*.
 - `strerror()` is more useful when you want to customize your own output format.

```
err = pthread_create(&ntid, NULL, thr_fn, "new thread: ");
if (err != 0) {
    fprintf(stderr, "can't create thread: %s\n", strerror(err));
    exit(1);
}
```

Pthread Synchronization

- **Semaphores** (#include `<semaphore.h>`)
 - Resources counter
 - Allows a limited number of threads in the critical section
- **Mutexes**
 - Binary semaphores
 - Allows only one thread in the critical section
- **Condition Variables**
 - Communicate information about the status of some shared data

Mutex

- Header file: `#include<pthread.h>`
- `pthread_mutex_init`
- `pthread_mutex_destroy`
- `pthread_mutex_lock`
- `pthread_mutex_unlock`

pthread_mutex_init

- Set the attribute of a mutex
- `int pthread_mutex_init(pthread_mutex_t *restrict mutex, const pthread_mutexattr_t *restrict attr);`
- We have three types of mutex
 - fast (default)
 - recursive
 - error check

pthread_mutex_lock

- If the mutex is unlocked: It becomes locked and owned by the calling thread
- If the mutex is locked by **another** thread, the calling thread is suspended until it is unlocked.
- What if the calling thread has already locked the mutex ?
 - pthread_mutex_lock(mutex);
 - pthread_mutex_lock(mutex);
 - ...

If the calling thread has already locked the mutex

- We have three types of mutex (statically defined below)
 - fast (default)
 - `pthread_mutex_t fast = PTHREAD_MUTEX_INITIALIZER`
 - Calling thread suspended on a locked mutex until it is unlocked (deadlock)
 - recursive
 - `pthread_mutex_t recursive = PTHREAD_RECURSIVE_MUTEX_INITIALIZER_NP`
 - Calling thread succeeds and returns immediately if a mutex is locked, recording the number of times it has locked the mutex. (Used when you unlock the mutex)
 - error check
 - `pthread_mutex_t errchk = PTHREAD_ERRORCHECK_MUTEX_INITIALIZER_NP`
 - Calling thread returns immediately if a mutex is locked
- Note: They are different only when the mutex is already locked by the calling thread.

pthread_mutex_trylock

- It behaves identically to pthread_mutex_lock, except that it does not block the calling thread if the mutex is already locked by **another** thread (or by the calling thread in the case of a "**fast**" mutex)
- It returns immediately with error code "EBUSY".

pthread_mutex_unlock

- Unlocks the mutex. The mutex is assumed to be locked and owned by the calling thread
- Be careful if the mutex is "recursive" or "error check" type. They will have different behaviors.
 - Refer to MAN page:
 - http://linux.die.net/man/3/pthread_mutex_lock

Condition Variable

- Header file: `#include<pthread.h>`
 - `pthread_cond_init`
 - `pthread_cond_wait`
 - `pthread_cond_broadcast`
 - `pthread_cond_signal`
-
- Note: Condition Variable will always be used with a mutex

pthread_cond_init

- Unlike mutex, the attribute of a condition variable is actually ignored. So "NULL" is always specified.
- Or it can be initialized statically,
 - `pthread_cond_t cond = PTHREAD_COND_INITIALIZER;`

pthread_cond_signal

pthread_cond_broadcast

- pthread_cond_signal
 - Restart one of the threads waiting on the condition
 - If there are more than one threads waiting on the same **cond**, it is not specified which thread will be restarted.
- pthread_cond_broadcast
 - Restart all the threads waiting on the condition variable

pthread_cond_wait

- Atomically unlocks the mutex (automatically)
- Waits for the cond to be signaled
- Thread is suspended and does not consume any CPU time until cond is signaled.
- Mutex must be locked by the calling thread before pthread_cond_wait()
- Before thread is recovered from the suspension at pthread_cond_wait(), it re-acquires (re-locks) mutex (automatically)

pthread_cond_destroy

- Test whether there are threads waiting on the condition variable

More helpful readings

- [POSIX thread \(pthread\) libraries](#)
- [Synchronizing Threads with POSIX Semaphores](#)
- [POSIX Threads Programming](#)
- [Multi-Threaded Programming With POSIX Threads](#)
- [Multithreaded Programming \(POSIX pthreads Tutorial\)](#)
- [Pthread Tutorial](#)