

Kmemleak

Kernel Memory Leak (kmemleak)

Contents

- [1 Introduction](#)
- [2 Prerequisites](#)
- [3 Usage](#)
- [4 References](#)

Introduction

Memory leak occurs when a piece of memory that was previously allocated by a programmer is not properly deallocated. **Kmemleak** is a feature of kernel, that helps in finding the memory leaks in **kernel space**. It only reports the list of memory leaks in the kernel space.

Valgrind is a similar tool to find the memory leak in **user space**.

Prerequisites

CONFIG_DEBUG_KMEMLEAK in kernel configuration need to be enabled as follows

1. Open configuration file.

```
$ make menuconfig
```

2. Enable CONFIG_DEBUG_KMEMLEAK option in Memory Debugging and the location is as below

Location:

- > Kernel Hacking
- > Memory Debugging
- > Maximum kmemleak early log entries (**change 400 to 8000**)

3. Once configuration is enabled, set the early log values in **Maximum kmemleak early log entries** to 4000. The default value of 400 may not work correctly in all configurations.

4. Build the kernel using following commands

```
$ make
$ make modules
$ make modules_install
$ make install
```

5. Reboot, once the build is complete.

6. Upon rebooting, mount debugfs file system.

```
$ mount -t debugfs nodev /sys/kernel/debug/
```

Note : If `/sys/kernel/debug` is already mounted, you will get response message `'/sys/kernel/debug/busy'`

7. Once `debugfs` is mounted, we can see a file **kmemleak** under `debugfs` mounted location.

```
$ cat /sys/kernel/debug/kmemleak
```

The above `/sys/kernel/debug/kmemleak` contains the information about any memory leak that has been detected.

Usage

1. Now let us see the usage of `kmemleak` to detect memory leak in the sample source (`sample_module.c`)

```
#include <linux/init.h>
#include <linux/module.h>
#include <linux/kernel.h>
#include <linux/vmalloc.h>
void myfunc(void)
{
    char *ptr;
    ptr = vmalloc(512);
}
static int my_init(void)
{
    printk(KERN_ALERT "Hello World");
    myfunc();
    return 0;
}
static void my_exit(void)
{
    printk(KERN_ALERT "Exit");
}

module_init(my_init);
module_exit(my_exit);

MODULE_LICENSE("GPL v2");
MODULE_AUTHOR("Sample_module");
```

2. Compile the above code and insert the `sample_module.ko` module.

```
$ insmod sample.ko
```

3. The memory leak detection thread runs periodically, now perform the test.

```
$ echo scan > /sys/kernel/debug/kmemleak
```

4. Let us check for leak.

```
$ cat /sys/kernel/debug/kmemleak
```

5. If leak detected, you should see a prompt like:

```
unreferenced object 0xffffc90000667000 (size 512):
  comm "insmod", pid 5602, jiffies 4297141725 (age 18.452s)
  hex dump (first 32 bytes):
    0a 00 00 00 0e 00 00 00 6b 56 00 00 00 00 00 00  .....kV.....
```

```

e8 82 00 00 00 00 00 00 0a 00 00 00 0e 00 00 00 .....
backtrace:
[<ffffffff817e449a>] kmemleak_alloc+0x4a/0xa0
[<ffffffff811ca55c>] __vmalloc_node_range+0x1ac/0x290
[<ffffffff811ca904>] vmalloc+0x54/0x60
[<fffffffa054e041>] my_init+0x21/0x30 [sample]
[<ffffffff81002190>] do_one_initcall+0x50/0x190
[<ffffffff81182b89>] do_init_module+0x60/0x1f1
[<ffffffff81107094>] load_module+0x21a4/0x2910
[<ffffffff811079f6>] SYSC_finit_module+0x96/0xd0
[<ffffffff81107a4e>] Sys_finit_module+0xe/0x10
[<ffffffff817eeb3b>] entry_SYSCALL_64_fastpath+0x1e/0xad
[<ffffffffffffffff>] 0xffffffffffffffff

```

The bold text above represents the leak detected in my_init function.

Note : The memory leak detector code may take sometime to detect the leak, so repeat step 3 and 4.

6. To clear the list of all current memory leaks

```
$ echo clear > /sys/kernel/debug/kmemleak
```

References